

IDUG Solutions Journal

IDUG Solutions Journal

Copyright © 2010 International DB2 Users Group

Table of Contents

From The Editor's Desk	1
Using DB2 pureXML and ODF Spreadsheets	2
Abstract	2
Introduction	2
The ODF specification	2
DB2 pureXML	2
The ODF document	2
A bit of hacking	4
Importing and XML document into a DB2 table	6
Converting the content.xml into row-column format	7
Getting the cell values from the XML file	7
The XPath expression	7
The Required Query	8
Shopping Prices	10
Building the List of Shopping Prices	10
XQuery Update to build content.xml document	12
A quick side note on XQuery Update	12
Generating the content.xml for the shopping prices	12
Building the <table:table-row> elements	13
Next steps	16
Resources	17
Author Bio	17
Giving You A Reason	18
Should you upgrade to DB2 10?	18
When and how should I upgrade to DB2 10?	18
DB2 9: Robust, Scalable, Available and Easily Manageable	19
DB2 10: Cut costs and improve performance	19
Questions for you	19
What version are you running?	20
DB2 and Storage Management	21
DB2 Storage Basics	21
Important DB2 for z/OS Storage Issues	23
A Little Bit About Modern Disk Arrays	25
Cache Versus Buffer	25
A Little Bit About DFSMS	26
What About Extents?	27
Another Thing to Think About	28
Best Practices	28
Summary	29
Author Bio	29
Don't Flip Out! How to Stop Your Query Access Plans from Flopping	30
A Reactive Approach	30
A Proactive Approach	31
Writing SQL Statements	31
Data Server Configuration	32
Configuration Parameters	33
Table Space Attributes	33
Optimization Class	34
Catalog Statistics	34
Column Group Statistics	35
LIKE Predicate Statistics	35

Statistical Views	36
Handling Exceptions	37
Access Plan Reuse	37
Optimization Profiles	38
Conclusion	40
.....	40

List of Figures

1. Contents of an ODF document file	3
2. Example ODF spreadsheet (row highlighted)	4
3. XML representation of ODF spreadsheet (row highlighted)	4
4. Edited content.xml showing added row	5
5. Spreadsheet Resulting from XML Document Edit	5
6. Mozilla Firefox with XPather Available	8
7. XPather Dialog Box	8
8. content.xml for an empty spreadsheet	13
9. ShopPrices.ods	15
10. Document Properties in OpenOffice	16
11. Sample meta.xml (partial)	16
1. DB2 Storage Groups vs SMS Storage Groups	22
2. DB2 Data Set Naming Convention	23
3. DBA vs Storage Administration	24

List of Tables

1. XML files in an ODF document file	3
2. Prefix and namespace URI used in content.xml document	8
1.	33
2.	33

List of Examples

1. DB2 table to hold content.xml file of the ODF spreadsheet	6
2. Contents of flat file to be used for loading DOCCONTENTXML table	6
3. IMPORT statement to load content.xml into DOCCONTENTXML table	6
4. XPath expression to get first row, second column value	7
5. Query to extract cell values from content.xml into tabular format	9
6. Result of SELECT on the view V_CONTENT	10
7. DDL for table to hold the shopping prices	10
8. View with list and cost of items with grand total	11
9. List of Items with Cost and Total Amount	12
10. A sample <table:table-row> element	13
11. Table to hold generated <table:table-row> elements	14
12. Insert generated <table:table-row> elements with <dummy> parent	14
13. Updated content.xml document	15
1. An example of an optimization profile	39

From The Editor's Desk

Philip Nelson

Welcome to the new look IDUG Solutions Journal, developed in response to your feedback received over the last few months.

Our main feature article this month is a case study in producing and consuming documents in ODF (Open Document Format), as used by the free OpenOffice.org suite, using DB2 pureXML. The techniques described in this article could be readily applied to a wide variety of other use cases.

We are also pleased to have a full complement of columns from our regular contributors. Roger Miller talks about the choices to be made when considering a DB2 for z/OS upgrade. John Hornibrook discusses how to guard against queries going bad. And Craig Mullins looks at the relationship which needs to be built between DBAs and Storage Managers. Next time we hope to add to our stable of regular columnists so watch this space.

As I said before this is the first issue of a new look ISJ. It is still very much a work in progress as we learn a bunch of new tools. This first issue is not exactly a "thing of beauty" but we hope this will get better as our skills build. At least I believe there is plenty of good technical content in it.

We are urgently in need of volunteers to help with the ISJ. These fall into three categories -

1. First and foremost we need authors. We'd be delighted to hear from anyone who would be willing to spend time working on a technical article for the magazine.
2. We also need reviewers to join the ERB (Editorial Review Board). Reviewers can be of two types : technical and grammatical. If you think you could contribute in either of these areas then please drop me a line. You don't need to be a cross-platform DB2 expert to be a good technical reviewer : indeed when I started with ISJ I learnt a lot as I checked the technical correctness of articles against the publicly available documentation.
3. Finally we need folks to help with the preparation of the content. If anyone out there has experience using the DocBook standard (or would be willing to learn) then we would be delighted to hear from you. Our aim is to turn the ISJ back into a publication not only with the best quality of content, but which also looks professional.

We hope to have the next issue of the ISJ available before the end of the year. Until then ...

... don't just do it, DB2 it !!!

Phil Nelson (ISJ Executive Editor)

Using DB2 pureXML and ODF Spreadsheets

Abstract

ODF (Open Document Format) is a file format for office documents. At a high level, the ODF specification requires five mandatory XML documents and other optional items. This article will describe how DB2 pureXML could be used to handle ODF spreadsheet documents. Other ODF documents such as word documents, presentations, formulas, etc. could also be handled by DB2 because the ODF specification refers mainly to an archive of XML documents. However, for the sake of this article, we are considering only spreadsheets because it is easier to co-relate the rows and columns of the spreadsheet document and that of a DB2 table. This article can even be further extended into an on-line document editing software with .ods as file format and DB2 pureXML as the database.

Introduction

The ODF specification and DB2 pureXML are introduced below.

The ODF specification

The Open Document Format specification was originally developed by Sun [<http://www.sun.com/>] and the standard was developed by OASIS Open Document Format for Office Applications (Open Document) TC – OASIS ODF TC. This standard is based on XML format originally created and implemented by the OpenOffice.org office suite. In addition to being a free and open OASIS standard, it is published (in one of its version 1.0 manifestations) as an ISO/IEC international standard, ISO/IEC 26300:2006 Open Document Format for Office Applications (OpenDocument) v1.0 [http://www.iso.org/iso/iso_catalogue/catalogue_tc/catalogue_detail.htm?csnumber=43485].

DB2 pureXML

DB2 pureXML [<http://www-01.ibm.com/software/data/db2/xml/>] is IBM software for management of XML data. It eliminates much of the work typically involved in the management of XML data. From validation of XML documents, to generation of documents and support for querying XML data in the form of XQuery [<http://www.w3.org/TR/xquery/>] and traditional SQL, DB2 pureXML is a complete solution to manage XML data in enterprise-scale databases.

The ODF document

An ODF document is a zipped file of five mandatory XML files and other optional components. This can be seen by running your favorite unzip utility to list the components. Here is the result after running the unzip command on an .ods file from command line to list the files. The XML files are highlighted below.

Figure 1. Contents of an ODF document file

```
nsubrahm@NSUBRAHM:~/Desktop/dw$ unzip -l Shop.ods
Archive:  Shop.ods
  Length      Date    Time    Name
-----
         46  05-13-09  09:44  mimetype
          0  05-13-09  09:44  Configurations2/statusbar/
          0  05-13-09  09:44  Configurations2/accelerator/current.xml
          0  05-13-09  09:44  Configurations2/floater/
          0  05-13-09  09:44  Configurations2/popupmenu/
          0  05-13-09  09:44  Configurations2/progressbar/
          0  05-13-09  09:44  Configurations2/menubar/
          0  05-13-09  09:44  Configurations2/toolbar/
          0  05-13-09  09:44  Configurations2/images/Bitmaps/
       7931  05-13-09  09:44  content.xml
       6482  05-13-09  09:44  styles.xml
        777  05-13-09  09:44  meta.xml
       5459  05-13-09  09:44  Thumbnails/thumbnail.png
       7228  05-13-09  09:44  settings.xml
       1896  05-13-09  09:44  META-INF/manifest.xml
-----
      29819                  15 files
```

Here is a short description of XML files as seen above.

Table 1. XML files in an ODF document file

File name	Usage	Location in zipped file
manifest.xml	Information about the files contained in the package	/META-INF
styles.xml	Styles used in the document content and automatic styles used in the styles themselves	/ (root)
settings.xml	Application-specific settings, such as the window size or printer information	/ (root)
content.xml	Document content and automatic styles used in the content	/ (root)
meta.xml	Document meta information, such as the author or the time of the last save action	/ (root)

For this article, we are interested in the content.xml file which will hold the contents of the document. The document used for this article is a simple shopping list named Shop.ods. Below is the spreadsheet as it appears in OpenOffice, with the first row highlighted.

Figure 2. Example ODF spreadsheet (row highlighted)

	A	B	C
1	Potato	2.0	kg
2	Onion	1.0	kg
3	Cucumber	0.5	kg
4	Capsicum	0.5	kg
5	Carrot	0.5	kg
6	Green chilli	0.25	kg
7	Tomato	0.25	kg
8	Curd	0.5	L
9	Banana	2	doz
10	Milk	2	L
11	Groundnut Oil	5	L

And below we have part of the equivalent content.xml file, with the XML markup for the row highlighted.

Figure 3. XML representation of ODF spreadsheet (row highlighted)

```
-<office:document-content office:version="1.2">
  <office:scripts/>
  +<office:font-face-decls></office:font-face-decls>
  +<office:automatic-styles></office:automatic-styles>
  -<office:body>
    -<office:spreadsheet>
      -<table:table table:name="Sheet1" table:style-name="ta1" table:print="false">
        <table:table-column table:style-name="co1" table:default-cell-style-name="Default"/>
        <table:table-column table:style-name="co2" table:default-cell-style-name="Default"/>
        <table:table-column table:style-name="co3" table:default-cell-style-name="Default"/>
        -<table:table-row table:style-name="ro1">
          -<table:table-cell office:value-type="string">
            <text:p>Potato</text:p>
          </table:table-cell>
          -<table:table-cell table:style-name="ce1" office:value-type="float" office:value="2">
            <text:p>2.0</text:p>
          </table:table-cell>
          -<table:table-cell office:value-type="string">
            <text:p>kg</text:p>
          </table:table-cell>
        </table:table-row>
      +<table:table-row table:style-name="ro1"></table:table-row>
```

A bit of hacking

Now that we know a bit about an ODF document, let us try to create one without actually running a standard software dealing with ODF documents. As noted above, the content.xml file holds the data entered while creating the spreadsheet. So, if this file is changed and zipped along with the rest of the files in the original archive (that is, the .ods file), we will be having a “legal” ODF spreadsheet. To test this, follow the steps below.

1. Open the content.xml file in a regular text editor
2. In the figure above, the element <table:table-row> has been highlighted. Type in the contents of the highlighted section immediately as a sibling of another <table:table-row> element.
3. In the pasted element, make a change in any of the table:table-cell/text:p element say, something like this -

Figure 4. Edited content.xml showing added row

```
-<table:table-row table:style-name="ro1">
-  <table:table-cell office:value-type="string">
    <text:p>Onion</text:p>
  </table:table-cell>
-  <table:table-cell table:style-name="ce1" office:value-type="float" office:value="1">
    <text:p>1.0</text:p>
  </table:table-cell>
-  <table:table-cell office:value-type="string">
    <text:p>kg</text:p>
  </table:table-cell>
</table:table-row>
-  <table:table-row table:style-name="ro1">
-    <table:table-cell office:value-type="string">
      <text:p>Coffee Powder</text:p>
    </table:table-cell>
-    <table:table-cell table:style-name="ce1" office:value-type="float" office:value="250">
      <text:p>250</text:p>
    </table:table-cell>
-    <table:table-cell office:value-type="string">
      <text:p>gms</text:p>
    </table:table-cell>
  </table:table-row>
```

4. Save the edited content.xml file in the same folder where the Shop.ods file was unzipped into, replacing the older one.
5. Select all the contents (that is, the ones mentioned in the manifest.xml file) of the unzipped folder and zip them. Name the zipped file as NewShop.ods.
6. Now, open NewShop.ods in OpenOffice or IBM Lotus Symphony and confirm that the hacking is successful. Note that OpenOffice successfully opens a file with a .zip extension, so long it is a valid ODF document whereas Symphony requires the extension to be .ods.

Figure 5. Spreadsheet Resulting from XML Document Edit

	A	B	C	D
1	Potato	2 kg		
2	Onion	1 kg		
3	Coffee Powder	250 gms		
4	Cucumber	0.5 kg		
5	Capsicum	0.5 kg		
6	Carrot	0.5 kg		
7	Green chilli	0.25 kg		
8	Tomato	0.25 kg		
9	Curd	0.5 L		
10	Banana	2 doz		
11	Milk	2 L		
12	Groundnut Oil	5 L		

It is now clear that, to “read” an existing ODF spreadsheet in its simplest form, we have to navigate the paths of the content.xml XML document. Similarly, to “modify” an existing ODF spreadsheet, we should

generate the content.xml XML document correctly. One can, of course, extend this to create and/or modify other XML documents in the .ods archive. But, for the sake of this article, we will concentrate only on the creation of content.xml XML document.

With this knowledge of XML documents in an ODF document archive, we will now proceed to demonstrate how to use DB2 pureXML with ODF documents. In the remainder of this article we will consider an ODF spreadsheet of a simple shopping list containing item names, units of measurement (UOM) and the number of units required. We will demonstrate how to -

1. Import the content.xml document of this ODF spreadsheet into an XML column within a DB2 table
2. Extract information from the content.xml document into relational format.
3. Merge data from a relational table with the content.xml document in the XML column to produce a new spreadsheet.

Importing and XML document into a DB2 table

We create a table to hold the XML document. As well as an XML column to hold the document it will also have an INTEGER column as a primary key which we will use to differentiate between multiple documents. The DDL for this table is as follows -

Example 1. DB2 table to hold content.xml file of the ODF spreadsheet

```
-- DOCCONTENTXML table to hold the content.xml file of a ODF document.
-- The table is keyed on DOCID a running sequence number.
CREATE TABLE ODF.DOCCONTENTXML
(
  DOCID          INTEGER NOT NULL,
  DOCCONTENT XML    NOT NULL
);
```

Inserting the XML document into the table could be done using a simple INSERT SQL statement. However since we have a considerably sized XML document, we will use the IMPORT statement from the command line. To use the IMPORT statement, first, save the content.xml in a folder as say, ODSFolder and create a flat file with a single record. Save the flat file as load.txt. This record will have two values separated by a comma as shown below -

Example 2. Contents of flat file to be used for loading DOCCONTENTXML table

```
1,<XDS FIL='content.xml' />
```

Example 3. IMPORT statement to load content.xml into DOCCONTENTXML table

```
IMPORT FROM load.txt OF DEL
XML FROM ODSFolder
INSERT INTO ODF.DOCCONTENTXML;
```

Alternatively, IBM Data Studio Developer can be used to load into XML column as it offers a convenient GUI.

Converting the content.xml into row-column format

Getting the cell values from the XML file

To get to the cell values of the spreadsheet, we will start with the required XPath expression. Then, we will develop a working query around the XPath expression to get the results.

The XPath expression

Going back to Table 2, the left column shows the first row of the spreadsheet highlighted. The right column is a part of the generated content.xml file and some of the elements collapsed. The relevant part are put in red boxes. The first box shows the table:table element with an attribute table:name whose value is "Sheet1". It is the name of the sheet in the complete spreadsheet. This element is repeated as many times as there are sheets in the spreadsheet with the table:name attribute set to the names as set by the user. The second box is the representation of the highlighted row in the left column. Thus, to get to, say, column B of the first row, the path to be navigated (or, the XPath expression) from the root of XML document would be as below.

Example 4. XPath expression to get first row, second column value

```
/office:document-content/office:body/office:spreadsheet/ (cont ...)  
  table:table[1]/table:table-row[1]/table:table-cell[2]/text:p
```

In other words, we start from the office:document element to office:body and then to office:spreadsheet to look for the first table:table element. The first element represents the first sheet of the spreadsheet document. Then, to get to the second column of first row we navigate to table:row[1]/table:cell[2] and finally to text:p to get the actual text of the cell.

With Firefox Mozilla browser, the above XPath expression can be easily tested using the add-on XPather (see Resources below). We follow the steps below :-

1. Open the content.xml in Firefox Mozilla browser
2. Right-click on the page and select 'Show in XPather'

Prefix	Namespace
table	"urn:oasis:names:tc:opendocument:xmlns:table:1.0"
text	"urn:oasis:names:tc:opendocument:xmlns:text:1.0"

The query shown below will use the XML namespaces and XPath expression that we produced in the previous sections to extract the cell values from the loaded content.xml file. The highlighted text shows the XPath expression used to extract cell values. The query has been wrapped into a view for ease of use -

Example 5. Query to extract cell values from content.xml into tabular format

```
CREATE VIEW ODF.V_CONTENT AS
(
  SELECT
  X.*
  FROM
  XMLTABLE(
    XMLNAMESPACES(
      'urn:oasis:names:tc:opendocument:xmlns:office:1.0' AS "office",
      'urn:oasis:names:tc:opendocument:xmlns:table:1.0' AS "table",
      'urn:oasis:names:tc:opendocument:xmlns:text:1.0' AS "text"
    ),
    'db2-fn:sqlquery(
      "SELECT
      DOCCONTENT
      FROM
      ODF.DOCCONTENTXML
      WHERE
      DOCID = 1"
    )
    /office:document-content/office:body/office:spreadsheet/
    table:table[1]/table:table-row'
  COLUMNS
    "ITEMNAME" CHARACTER (20) PATH 'table:table-cell[1]/text:p',
    "QUANTITY" DECIMAL (4,2) PATH 'table:table-cell[2]/text:p',
    "UOM" CHARACTER (5) PATH 'table:table-cell[3]/text:p'
  ) AS X
);
```

A simple SELECT query on this view shows that we have successfully 'converted' the content.xml document in the ODF archive for spreadsheet into a row-column format.

Example 6. Result of SELECT on the view V_CONTENT

```
SELECT * FROM ODF.V_CONTENT ;
```

ITEMNAME	QUANTITY	UOM
Potato	2.00	kg
Onion	1.00	kg
Cucumber	0.50	kg
Capsicum	0.50	kg
Carrot	0.50	kg
Green chilli	0.25	kg
Tomato	0.25	kg
Curd	0.50	L
Banana	2.00	doz
Milk	2.00	L
Groundnut Oil	5.00	L

11 record(s) selected.

Shopping Prices

We will now build the content.xml for an ODF spreadsheet that will hold the shopping list – discussed above – with an additional column for cost of items in the shopping list. We will also provide an additional row for the grand total of the items in shopping list.

Building the List of Shopping Prices

We create a simple table that holds the unit price of shopping items with the DDL shown below.

Example 7. DDL for table to hold the shopping prices

```
--Table to hold the unit price of shopping items
CREATE TABLE ODF.SHOPPER
( ITEMNAME    CHARACTER(50),
  UNIT_PRICE  DECIMAL (4,2)
);
```

This table is inserted with values such that all the item names are available in the SHOPPER table. To obtain the list of item names and the cost of items, we use the following query. Note that, this query uses the view V_CONTENT. We could have used the base table directly. Usage of the view makes the query 'compact' and easier to follow.

Example 8. View with list and cost of items with grand total

```
CREATE VIEW ODF.V_SHOP_COST AS
(
  SELECT
    COALESCE(P.ITEMNAME, ' TOTAL')
    ITEMNAME,
    SUM(P.QUANTITY) QTY,
    COALESCE(P.UNIT_PRICE, 0) UNIT_PRICE,
    SUM(P.COST) COST
  FROM
    (
      SELECT
        A.ITEMNAME,
        A.QUANTITY,
        B.UNIT_PRICE,
        A.QUANTITY*B.UNIT_PRICE COST
      FROM
        ODF.V_CONTENT A,
        ODF.SHOPPER B
      WHERE
        UPPER(A.ITEMNAME) = B.ITEMNAME
    ) P
  GROUP BY
    ROLLUP(P.ITEMNAME, P.QUANTITY, P.UNIT_PRICE)
  HAVING
    (P.ITEMNAME IS NULL OR
     (P.ITEMNAME      IS NOT NULL AND
      P.QUANTITY      IS NOT NULL AND
      P.UNIT_PRICE     IS NOT NULL AND
      SUM(P.COST)      IS NOT NULL AND
      SUM(P.QUANTITY) IS NOT NULL
    )
  )
);
```

The result of querying this view, ordering the data by item name (descending) is shown below.

Example 9. List of Items with Cost and Total Amount

ITEMNAME	QTY	UNIT_PRICE	COST
-----	-----	-----	-----
Tomato	0.25	6.00	1.50
Potato	2.00	5.50	11.00
Onion	1.00	4.00	4.00
Milk	2.00	40.00	80.00
Groundnut Oil	5.00	30.00	150.00
Green chilli	0.25	8.00	2.00
Curd	0.50	20.00	10.00
Cucumber	0.50	10.00	5.00
Carrot	0.50	15.00	7.50
Capsicum	0.50	20.00	10.00
Banana	2.00	6.00	12.00
TOTAL	14.50	0.00	293.00

XQuery Update to build content.xml document

At this point, we have 'shredded' the shopping list (from the content.xml of ODF spreadsheet), joined with the relational table having the unit prices of items and generated a relational output of the cost of items in the shopping list with the grand total. Our aim is now to convert this relational output into the format of content.xml so that the result is a complete ODF spreadsheet.

One way of building the content.xml is to start from scratch and building the whole of the document. This approach would be quite tedious. A better way would be to generate the relevant part in the content.xml part and then insert into an otherwise empty content.xml document. This approach will also allow us to make use of the XQuery Update facility in DB2 pureXML.

A quick side note on XQuery Update

XQuery is a query and functional programming language designed to query XML documents. It became a W3C Candidate Recommendation on 23rd January, 2007. XQuery Update facility is an extension to XQuery that allows update (insert, modify and delete) of the XML documents. It became a W3C Candidate Recommendation on 14th March, 2008.

Generating the content.xml for the shopping prices

To start with, we will create a dummy content.xml that is complete and ODF compliant in all respects except it would resemble an empty spreadsheet. Next, we will generate XML sequences that are correct for a single row within the spreadsheet. Then, we will collect all of these XML sequences and insert into the dummy content.xml document.

To create the dummy document, we simply open the spreadsheet for the shopping list, delete all rows and columns and save it with another name. Then, we unzip this document and extract the content.xml document. Here is how this document would look -

Figure 8. content.xml for an empty spreadsheet

```
-<office:document-content office:version="1.2">
  <office:scripts/>
  +<office:font-face-decls></office:font-face-decls>
  +<office:automatic-styles></office:automatic-styles>
  -<office:body>
    -<office:spreadsheet>
      -<table:table table:name="1" table:style-name="ta1" table:print="false">
        <table:table-column table:style-name="co1" table:default-cell-style-name="ce1"/>
        <table:table-column table:style-name="co2" table:default-cell-style-name="ce1"/>
        <table:table-column table:style-name="co3" table:default-cell-style-name="Default"/>
        <table:table-column table:style-name="co4" table:default-cell-style-name="Default"/>
        <table:table-column table:style-name="co5" table:default-cell-style-name="Default"/>
      </table:table>
    </office:spreadsheet>
  </office:body>
</office:document-content>
```

Insert this dummy document into the ODF.DOCCONTENTXML table with DOCID as zero using the IMPORT method described before. Our aim now is to update this document such that, `<table:table-row>` elements are added as children of `<table:table>` and as siblings of `<table:table-column>`.

Building the `<table:table-row>` elements

In our example, each row of the spreadsheet would have four columns. Therefore, each row element, as required in content.xml, would have four children (one per column). A sample of `<table:table-row>` is shown below. The XML namespace declarations have been omitted for the sake of brevity.

Example 10. A sample `<table:table-row>` element

```
<table:table-row table:style-name="ro1">
  <table:table-cell office:value-type="string">
    <text:p>Tomato</text:p>
  </table:table-cell>
  <table:table-cell office:value-type="string">
    <text:p>.25</text:p>
  </table:table-cell>
  <table:table-cell office:value-type="string">
    <text:p>6.00</text:p>
  </table:table-cell>
  <table:table-cell office:value-type="string">
    <text:p>1.5000</text:p>
  </table:table-cell>
</table:table-row>
```

To build the `<table:table-row>` elements we need to generate `<table:table-cell>` for each column of the row in the relational table and then aggregate them to a single `<table:table-row>`

element. We will use the XMLAGG function to build the elements and store in a table the generated elements so that it is easier to use while doing the update of XML document. The table to hold the generated elements is defined as shown below.

Example 11. Table to hold generated <table:table-row> elements

```
CREATE TABLE ODF.SHOP_COST_XML
(SHOP_COST XML);
```

Example 12. Insert generated <table:table-row> elements with <dummy> parent

```
INSERT INTO ODF.SHOP_COST_XML
(SHOP_COST)
(SELECT
XMLDOCUMENT(
  XMLELEMENT(NAME "dummy",
    XMLNAMESPACES(
      'urn:oasis:names:tc:opendocument:xmlns:office:1.0' AS "office",
      'urn:oasis:names:tc:opendocument:xmlns:table:1.0' AS "table",
      'urn:oasis:names:tc:opendocument:xmlns:text:1.0' AS "text"
    ),
    XMLAGG(
      XMLELEMENT(NAME "table:table-row",
        XMLATTRIBUTES('rol' AS "table:style-name"),
        XMLELEMENT(NAME "table:table-cell",
          XMLATTRIBUTES('string' AS "office:value-type"),
          XMLELEMENT(NAME "text:p", STRIP(O.ITEMNAME))
        ),
        XMLELEMENT(NAME "table:table-cell",
          XMLATTRIBUTES('string' AS "office:value-type"),
          XMLELEMENT(NAME "text:p", O.QTY )
        ),
        XMLELEMENT(NAME "table:table-cell",
          XMLATTRIBUTES('string' AS "office:value-type"),
          XMLELEMENT(NAME "text:p", O.UNIT_PRICE )
        ),
        XMLELEMENT(NAME "table:table-cell",
          XMLATTRIBUTES('string' AS "office:value-type"),
          XMLELEMENT(NAME "text:p", O.COST)
        )
      )
    )
  )
  ORDER BY O.ITEMNAME DESC
);
```


Finally, we are now ready to update our dummy XML document. The query to do the update is shown below. Note that, this query is not inserting the resulting XML document into any table. It is left to the reader to route the output to a table column, file, message queue, etc.

Example 13. Updated content.xml document

```
VALUES XMLQUERY(
'declare namespace
    office="urn:oasis:names:tc:opendocument:xmlns:office:1.0";
declare namespace
    table="urn:oasis:names:tc:opendocument:xmlns:table:1.0" ;
transform
copy $dummy := db2-fn:sqlquery("
SELECT DOCCONTENT FROM ODF.DOCCONTENTXML WHERE DOCID = 0
") ,
$rows := db2-fn:sqlquery("
SELECT * FROM ODF.SHOP_COST_XML
")
modify do
insert $rows/dummy/* as last
into
$dummy/office:document-content/office:body/office:spreadsheet/table:table
return $dummy
');
```

To test that our update has worked successfully run, we route the output of query above to a file called content.xml. Replace the content.xml file in the unzipped archive of our original shopping list spreadsheet with the one created now. Zip the files back and save it with name as ShopPrices.ods and open with OpenOffice. Here is the screenshot.

Figure 9. ShopPrices.ods



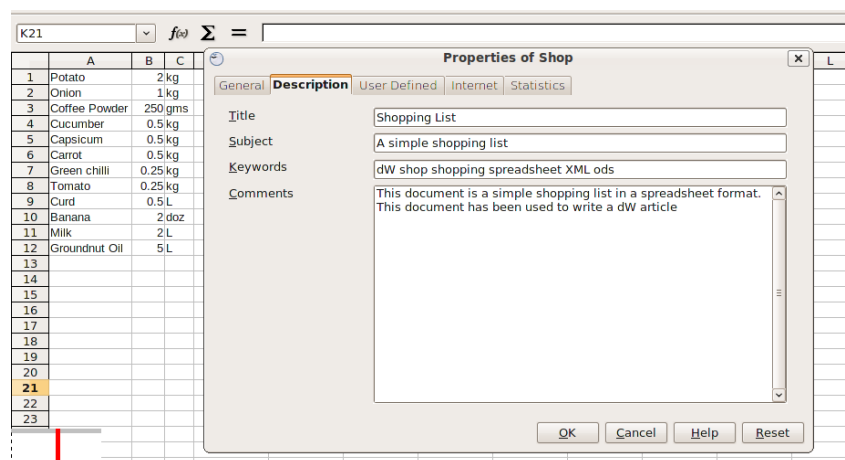
	A	B	C	D	E
1	Tomato	0.25	0.25	0.25	
2	Potato	2	2	2	
3	Onion	1	1	1	
4	Milk	2	2	2	
5	Groundnut Oil	5	5	5	
6	Green chilli	0.25	0.25	0.25	
7	Curd	0.5	0.5	0.5	
8	Cucumber	0.5	0.5	0.5	
9	Carrot	0.5	0.5	0.5	
10	Capsicum	0.5	0.5	0.5	
11	Banana	2	2	2	
12	TOTAL	14.5	14.5	14.5	
13					

And, we are done !

Next steps

Further enhancements may involve manipulating other XML documents in a ODF compliant archive. For example, the meta.xml document holds information about the document itself. This information can be accessed and set by going to File -> Properties... dialog box in OpenOffice. For example, for the document in this article (Shop.ods), this is how the dialog box looks like.

Figure 10. Document Properties in OpenOffice



The corresponding meta.xml document (partial snapshot) is shown below.

Figure 11. Sample meta.xml (partial)

```
--<office:meta>
  <meta:creation-date>2009-05-29T17:39:45</meta:creation-date>
  <meta:editing-duration>PT00H01M16S</meta:editing-duration>
  <meta:editing-cycles>2</meta:editing-cycles>
  <meta:generator>
    OpenOffice.org/3.0$Linux OpenOffice.org_project/300m15$Build-9379
  </meta:generator>
  <meta:initial-creator>Nagesh Subrahmanyam</meta:initial-creator>
  <dc:description>
    This document is a simple shopping list in a spreadsheet format. This document has been used to write a dW article
  </dc:description>
  <meta:keyword>dW shop shopping spreadsheet XML ods</meta:keyword>
  <dc:subject>A simple shopping list</dc:subject>
  <dc:title>Shopping List</dc:title>
  <dc:date>2009-05-29T17:39:47</dc:date>
  <dc:creator>Nagesh Subrahmanyam</dc:creator>
  <meta:document-statistic meta:table-count="3" meta:cell-count="36" meta:object-count="0"/>
  <meta:user-defined meta:name="Info 1"/>
  <meta:user-defined meta:name="Info 2"/>
  <meta:user-defined meta:name="Info 3"/>
  <meta:user-defined meta:name="Info 4"/>
</office:meta>
```

Thus, while generating the spreadsheet, one can write queries to generate the meta.xml document with correct values for creator, creation time-stamp, etc.

Resources

Important resources :

DB2 Express-C [http://www-01.ibm.com/software/data/db2/express/download.html]	– Version of DB2 which is free to download, develop, deploy and distribute.
DB2 v9 pureXML [http://www.ibm.com/db2/9]	All information about DB2 V9 pureXML
DB2 v9 Product Manuals [http://www.ibm.com/support/docview.wss?rs=71&uid=swg27009552]	Especially, refer the ones for XQuery and XML.
IBM Data Studio [http://www-01.ibm.com/software/data/studio/]	Govern, design, develop and deploy databases and data driven applications.
ODF Specification [http://www.oasis-open.org/committees/office/]	Link to the complete specification
OpenOffice.org XML essentials [http://books.evc-cit.info/book.php]	– Nice explanation of XML data formats.
Thinking XML : The open office file format [http://www.ibm.com/developerworks/xml/library/x-think15/]	An earlier developerWorks article discussing the file formats of OpenOffice.
Manage ODF and Microsoft Office 2007 documents with DB2 9 pureXML [http://www.ibm.com/developerworks/db2/library/techarticle/dm-0705gruber/]	An earlier developerWorks article on using ODF and Microsoft Office documents with DB2 9 pureXML.
XPather add-on [https://addons.mozilla.org/addon/1192]	Nifty tool for Firefox Mozilla users to quickly check an XPath expression
Update XML in DB2 9.5 [http://www.ibm.com/developerworks/db2/library/techarticle/dm-0710nicola/%5C]	Nice article on developerWorks explaining update of XML documents

Author Bio

Nagesh Subrahmanyam is an Advisory System Analyst with IBM India Pvt Ltd. He has over 9 years of experience with IBM System z technologies and is a keen fan of OSS technologies (Java, XML, etc.) especially on IBM System z. His other publications include the WebSphere MQ SupportPac MA1M, IBM Redbook 'XML Processing Options on z/OS' and a developerWorks article 'Submit batch jobs from Java on z/OS'. He can be reached at nsubrahm@in.ibm.com

Giving You A Reason

Compelling Reasons to Upgrade DB2

Roger Miller



Is your current bowl getting tight? What is limiting you? Is it CPU? Virtual storage? Latching? Catalog and directory? Utilities? Are you currently running DB2 9? V8? V7? Upgrading to a bigger bowl may be just what you need.

Should you upgrade to DB2 10?

To 10, or not to 10, when and how are the questions. Whether 'tis nobler in the mind to suffer the slings and arrows of outrageous limits, Or to take arms against a sea of troubles, And by upgrading, end them? To die: to sleep; No more; and by a sleep to say we end the heart-aches and the thousand natural shocks that old versions are heir to. 'tis a consummation devoutly to be wished. To die, to sleep. To sleep: perchance to dream: ay, there's the rub; For in that sleep of death what dreams may come when versions have shuffled off this mortal coil, must give us pause: There's the respect that makes calamity of too long life for old versions. [With abject apologies to the Bard and to Hamlet act 3 scene 1]

The answer to upgrading to 10 is a definite Yes. The question is not so much whether to upgrade as when and how to upgrade. If you are running DB2 9 today, then DB2 10 is in your near future, giving you more room to grow, with higher limits, lower costs, and more for less. If you are running DB2 V8 today, then you have a choice of jumping to DB2 9 or directly to DB2 10. So the key question is ...

When and how should I upgrade to DB2 10?

As of early August 2010, DB2 10 is in beta. Some of the key information for making this decision is known, but some is not yet. DB2 for z/OS V8 end of service is set for April 2012, 21 months from now. The unknown information includes the date for DB2 10 general availability, V8 extended service, and pricing, which will come in later announcements.

While DB2 10 is expected to be better than prior versions, it will have maturity, stability, and service delivery similar to other software and versions, with more defects at first, then fewer as the software matures. Determining when the software is ready for a specific customer and when the customer is ready for the software depends upon the specific customer resources, prior experience, and the value for the improvements versus the need for stability. Many customers depend upon tools or other software, and having that software that works with DB2 is a prerequisite. When this information is known, we can help answer the question. This web page can help. <http://www.ibm.com/support/docview.wss?uid=swg21006951> Content of the two versions is available, although the details of DB2 10 are still not public. Here is a summary of the two versions -

DB2 9: Robust, Scalable, Available and Easily Manageable

DB2 9 delivers CPU reductions for utilities that are generally in the range of 20% to 30%. Customers report saving terabytes of disk space using index compression. More CPU time is shifted to zIIP processors, reducing costs. Security is improved with more flexible trusted contexts and roles. Resilience is improved as more changes can be made while applications keep running. One table can be replaced quickly with a clone. Indexes and columns can be renamed. Many more utilities can be online.

DB2 9 delivers seamless integration of XML and relational data with pureXML and makes big strides in SQL for productivity and portability of applications. A new storage structure is introduced for large tables. Today's complex applications include both transactions and reporting, so performing both well is required. The key improvements for reporting are optimization enhancements to improve query and reporting performance and ease of use. More queries can be expressed in SQL with new SQL enhancements. Improved data is provided for the optimizer, with improved algorithms. Improved CPU and elapsed times can be achieved with the FETCH FIRST clause specified on a subquery. The INTERSECT and EXCEPT clauses make SQL easier to write.

DB2 10: Cut costs and improve performance

DB2 10 for z/OS provides the best reduction in CPU for transactions and batch in over 20 years. We expect most customers to reduce CPU times between 5% and 10% initially, with the opportunity for much more. Applications which can take advantage of additional benefits, such as hash access, can have larger CPU and memory reductions.

Scalability is the second major benefit, with the ability to run five to ten times as many threads in a single subsystem by moving 80% to 90% of the virtual storage above the bar. Schema evolution or data definition on demand enhancements improves availability. SQL and pureXML improvements extend usability and application portability for this platform. Productivity improvements for application developers and for database administrators are very important as data grows in scale and complexity.

Questions for you

The right answer is not "One size fits all." If we know the key factors for you, we can help you make a better choice. Here are some of the key objectives. Which ones are most important for you?

- Performance improves in both DB2 9 and 10, with larger CPU reductions in DB2 10.
- Scalability is improved a little in DB2 9 and a lot in DB2 10.
- Availability is enhanced in both, with more online changes in both.

- Security is made stronger and more flexible with roles in DB2 9 and with more options in DB2 10.
- Productivity is helped in both releases, with more improvements in DB2 10. Upgrading to DB2 9 is easier than to DB2 10 from DB2 V8.
- Stability is better in more mature versions.
- Skills: What skill set is available within your organization? Do you have people with the right skills and time to plan and run a project? DB2 planning workshops can help with education. Transition classes provide more education for one or both versions. Services could be used if the skills are not presently available.
- Technology adoption model: Are you using the latest technology that is being shipped, or are the operating system and hardware back-level? Is the technology one level back, two or more? This question tells both how much work will be required and the comfort level of your organization for running the latest version.
- Platform management practices: What are your platform management practices? What type of change management practices are in place? How robust is your testing for new software? What is the inventory of software and tools? How many vendors are involved? Which ones? Almost every vendor has software ready for DB2 9, but DB2 10 may take some time after general availability.
- Numbers of servers: How many LPARs and subsystems does the organization have? An organization that has 100 subsystems has a different set of challenges than does one that has 5 subsystems.
- Organizational considerations: What additional organizational factors, such as politics and policies, must be considered?

What version are you running?

Here are the primary recommendations for customers who are running various DB2 versions.

- DB2 9 If you are on DB2 9 today, then you are a good candidate for an early upgrade to DB2 10, especially if your custom is to move in the first year after general availability. Listen to reports from early customers and upgrade for the value.
- V8 If you are on DB2 V8 today, then the next questions are on timing for you and readiness for the new version. How soon after general availability do you normally upgrade? Are you still in the process of moving to NFM or have you recently finished V8 upgrade? If you just finished, then you probably will wait a few years and use the skip. If you have resources for an upgrade, but DB2 10 is too new, then DB2 9 is probably your next move. If you have the resources and can work with a new version, then skipping to DB2 10 may work for you.
- V7 If you are currently on DB2 V7, then upgrade to DB2 V8. Then you can use the skip version upgrade to DB2 10 in a few years.

DB2 has several new versions and upgrade paths for you to consider. This story will be changing, but you can hear the latest at IDUG conferences and on the web. DB2 9 is ready for you now. DB2 10 is still in beta, but is delivering higher limits, lower costs, and more for less.

- <http://www.ibm.com/software/data/db2/zos/db2-10/>
- <http://www.ibm.com/data/db2/zos>

DB2 and Storage Management

Craig S. Mullins

As an IT professional who uses DB2, you know that all database management systems rely on some form of persistent storage to maintain data. That means that the DBMS interoperates with operating system files, or data sets. As such, some form of storage management should be a key part of the database operations required of a DBA. Typically database storage means disk devices or subsystems, but it can also mean solid state disk, removable storage, or even trusty “old” tape devices.

DB2 Storage Basics

At the most basic level it is important to know that DB2 stores its data in VSAM Linear Data Sets (LDS). Each table space and index space you create requires at least one, possibly more, underlying VSAM data sets. DB2 uses VSAM Media Manager for its I/O operations. For every I/O, VSAM Media Manager builds a channel program and sends a request to the I/O supervisor.

But let’s back up a moment. The following items are the core of the storage-related objects and items you will need to know about for DB2 for z/OS:-

- DB2 Storage Groups : list of disk volumes
 - DB2 can also work with SMS so you will need to differentiate between DB2 storage groups and SMS storage groups (see Figure 1).
- Table Spaces : stored on disk as at least one VSAM LDS data set
- Indexes : stored on disk (in an index space) as at least one VSAM LDS data set
- System Data Sets
 - Active Log : stored on disk
 - Archive Logs : stored on disk or tape
 - BSDS : stored on disk
- Image Copy Backups : stored on disk or tape
- Other "stuff"
 - DB2 library data sets
 - Temporary data sets (used by utilities)

Figure 1. DB2 Storage Groups vs SMS Storage Groups

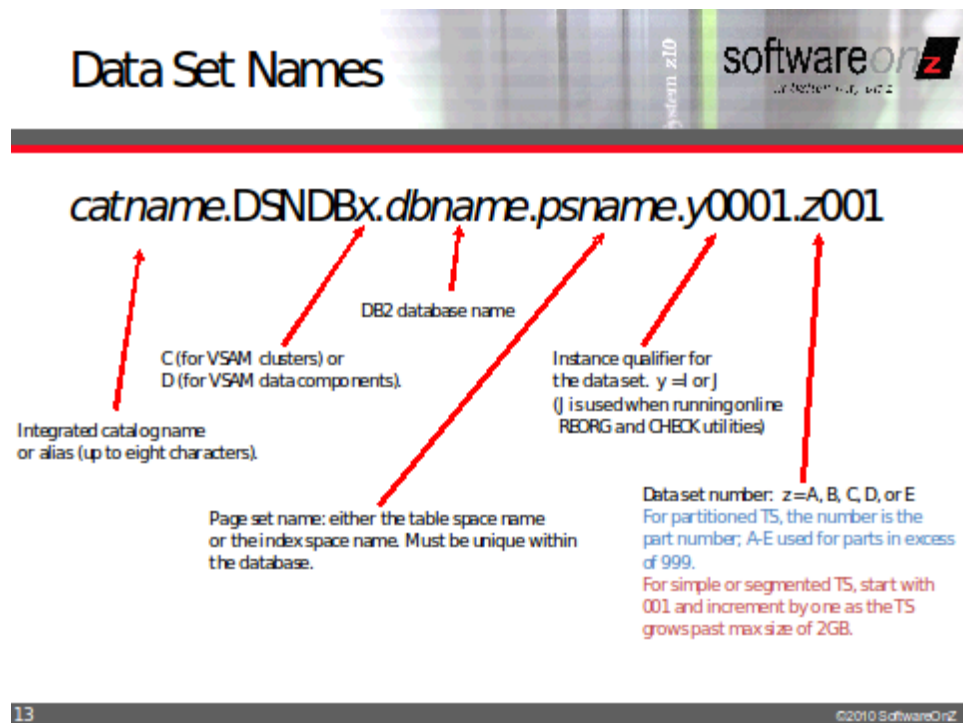
Storage Groups	
<ul style="list-style-type: none"> • DB2 Storage Group <ul style="list-style-type: none"> – DB2 uses STOGROUPs for disk allocation of the table and index spaces. – It is basically just a list of storage volumes <ul style="list-style-type: none"> • Or '*' for SMS – A STOGROUP can be assigned to: <ul style="list-style-type: none"> • Databases • Table spaces • Indexes 	<ul style="list-style-type: none"> • SMS Storage Group <ul style="list-style-type: none"> – An SMS Storage Group is a pool of disk volumes upon which SMS-managed data sets are placed. – A data set is placed on an appropriate volume within a Storage Group depending upon the Storage Class, Volume, Storage Group status, and available free space.

So you can see that there are multiple areas within DB2 that require storage and need to be managed. This article will touch upon most of these areas. But back to data set basics for the time being.

You may have noticed that I said that multiple data sets may be required so when does DB2 utilize multiple VSAM data sets for a table space or index? There are three situations where this will arise:-

1. When the object is partitioned, each partition will reside in a separate data set.
2. When a data set in a segmented or simple table space reaches its maximum size of 2 GB, DB2 can automatically create a new data set.
3. When the table space is cloned; each clone has its own underlying data set(s).

To understand how DB2 accommodates these situations we will need to take a look at how DB2 data sets are named. Figure 2 shows the naming convention for DB2 data sets. Although many of you reading this article may be familiar with this naming convention, a quick review is still a good idea. The database name and the page set name (table space or index space name) is part of the data set name. This is one of the reasons that you cannot have more than one table space or index space of the same name in the same database. The “interesting” part of the naming convention (if a naming convention can be interesting at all) comes at the end. We have an instance qualifier and a data set number.

Figure 2. DB2 Data Set Naming Convention

The instance qualifier is used when running online REORG and CHECK utilities. For an online utility DB2 uses a shadow data set and will switch from the current to the shadow after running the utility. So DB2 will switch the instance qualifier between I and J when you run online REORG and CHECK utilities. The numbers after the instance qualifier can change if you use clones. Although this is not the place for a comprehensive discussion of cloning let's skim the surface to understand what happens to this portion of the data set name. Basically, cloning creates a table with the exact same attributes as a table that already exists, except that it has no data. The clone is created using the ALTER TABLE SQL statement with the ADD CLONE parameter and the clone table is created in the same table space as the existing table... except in a different VSAM data set. The base table starts with I0001 in the data set name; the clone will be I0002. This can change because the SQL EXCHANGE statement flips the VSAM data sets.

Next we have the data set number, which appropriately enough, is used when a page set requires multiple data sets. The z is usually an "A", but it can be A, B, C, D, or E. For partitioned table spaces, the number is the partition number and A-E is used for partitions in excess of 999. So partition 1 would be A001 and partition 1,000 would be B001, and so on. For simple or segmented table spaces, the data set number starts at 001 and is incremented by one as the page set grows past the maximum size of 2GB.

These are the most basic storage "things" that you will need to know as you manage DB2 for z/OS storage.

Important DB2 for z/OS Storage Issues

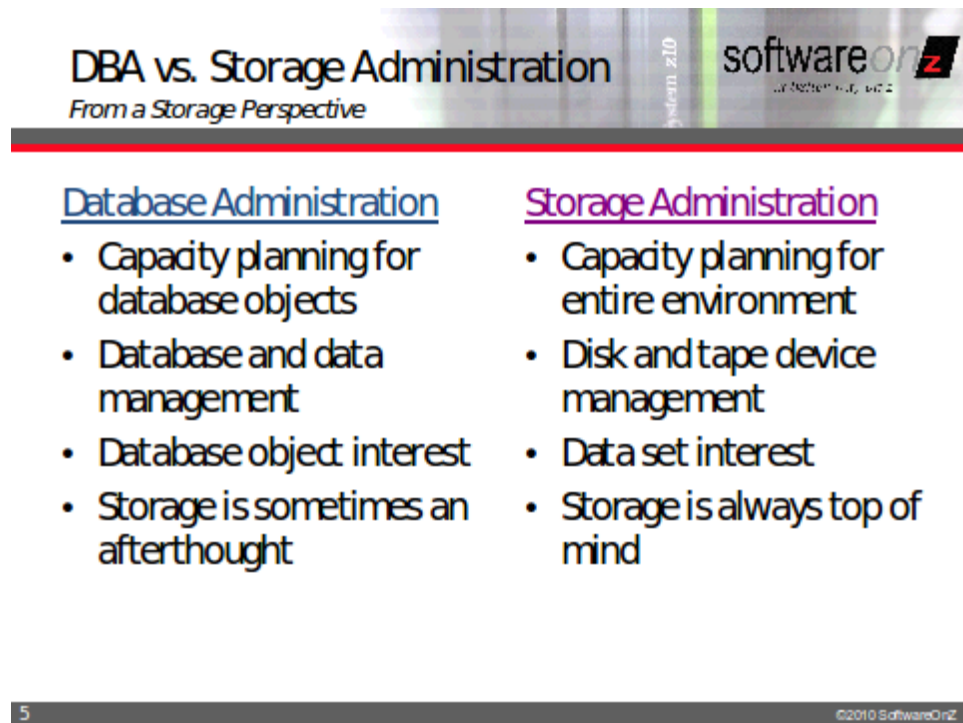
Although storage management can be an afterthought for the DBA it really shouldn't be. According to Gartner, Inc. the cost of managing storage is 4 to 10 times the initial cost of storage acquisition. And the growth rate for disk storage was 37% for the years 1996 through 2007. So storage issues are vitally important and unless it is managed appropriately it can be very costly. And unmanaged DB2 storage can result in system outages, which is the last thing any DBA wants to have happen, isn't it?

Even so, it is common for storage-related issues to be relegated to the backburner by DBAs. Let's face it, most robust mainframe organizations have an entire unit dedicated to storage management and adminis-

tration. And the DBA has enough to contend with without adding storage tasks to the list. But DBAs and storage administrators are concerned about different things – and that is the way it should be.

Refer to Figure 3. The DBA has a database-centric focus, whereas the storage administrator will focus on storage issues for the entire shop, focusing on devices and data sets. But the storage folks are not DB2 experts, nor should they be. Likewise, most DBAs are not storage experts. Making matters worse is that these two groups rarely communicate well.

Figure 3. DBA vs Storage Administration



<u>Database Administration</u>	<u>Storage Administration</u>
• Capacity planning for database objects	• Capacity planning for entire environment
• Database and data management	• Disk and tape device management
• Database object interest	• Data set interest
• Storage is sometimes an afterthought	• Storage is always top of mind

5 ©2010 SoftwareOnZ

So there is a gap between Database Administration, Storage Administration and Capacity Planning. Information is available to DBAs from various sources including RUNSTATS, STOSPACE, real-time statistics (RTS), DB2 Catalog, and so on, but the details are scattered all over the place and it can be difficult to gain a complete, accurate, and up-to-date picture. And any historical view into DB2 storage usage has to be managed manually.

Think about your environment for a moment and then reflect on whether or not you can easily answer the following questions:

- Do all of my databases have sufficient allocation to satisfy business requirements?
- Why is DB2 storage growing when our business is not?
- Am I wasting any storage?
- When will more storage be required?
- How much additional storage is needed?
- What needs to be done to align the additional storage with the DBMS?

Without a tool to capture, integrate, and manage information about your DB2 storage infrastructure answering these questions can be quite difficult.

A Little Bit About Modern Disk Arrays

Mainframe disk, or DASD, is usually equated to a 3380 or 3390. In other words, people think of physical hardware devices with a one-to-one relationship between a disk drive and a volume. The logical view is broken down as:

- Track size, or the number of bytes per track.
- Capacity, or the size of the device, in terms of number of tracks or gigabytes.
- Device address, which is a thread onto which I/O operations are serialized by the operating system

But the physical world is not the same as the logical anymore. Today's modern storage architecture uses disk arrays, or RAID. RAID stands for Redundant Array of Independent Disk; an array is the combination of two or more physical disk storage devices in a single logical device or multiple logical devices. The array is perceived by the system to be a single disk device. RAID can offer big benefits in terms of availability because the disks are typically hot-swappable, meaning that a drive can be replaced while the array is up and running. There are many levels of RAID technology, each delivering different levels of fault-tolerance and performance. A nice educational overview of the various levels of RAID is offered by Advanced Computer & Network Corporation at http://www.acnc.com/04_00.html. But what about mainframe disk arrays?

The RAMAC Virtual Array (RVA) came first for the mainframe and it was based on virtual disks that emulated 3380s and 3390s. The RVA dynamically maps functional volumes to physical drives. This mapping structure is contained in a series of tables stored in the RVA control unit. RVA was OEM'd from Storage Technology Corp. (STK) which is now part of Oracle.

The ESS (Enterprise Storage System), also known as Shark, followed when the STK OEM agreement expired. The ESS is scalable from 420GB to 55.9 TB. It offers improved performance over RAMAC, especially for prefetch and analytical queries.

And the latest and greatest IBM mainframe disk technology is the DS8000, which employs virtualized disk. It adds functionality for storage pool striping, thin provisioning, and quick initialization, among other innovations. Its capacity scales linearly from 1.1 TB up to 192 TB (up to 320 TB with turbo models).

Of course, IBM is not the only game in town for mainframe storage. EMC, Hewlett Packard, Hitachi, and Sun also offer modern disk arrays for the mainframe.

Cache Versus Buffer

Modern disk arrays cache data, but so does DB2. In the old days we used to disable disk cache for DB2, but no longer. There are two DSNZPARMs you should be aware of that can be setup to properly control disk caching for DB2 data sets.

First up is the SEQCACH parameter. The original meaning of this parameter, for 3390 DASD, was whether DB2 I/O should bypass the disk cache, but the meaning is different now because you do not want to bypass the cache on modern storage arrays. There are two options:

- BYPASS : the disk will perform Sequential Detection
- SEQ : creates an explicit Prefetch request; the recommendation is to use this setting for improved performance.

The second DSNZPARM is SEQPRES, which is similar to SEQCACH, but for DB2 LOAD and REORG utilities. If set to YES the Cache is more likely to retain pages for subsequent update, particularly when processing NPIs. The general recommendation is to set to SEQCACH to YES.

Keep in mind, too, that your storage administrators should be aware that DB2 buffering may cause DB2 data to use the disk cache differently than other non-database data sets. But that doesn't mean that DB2 is not benefiting from disk caching.

A Little Bit About DFSMS

DFSMS is IBM's Data Facility Storage Management System. It offers data management, backup and HSM software from IBM mainframes, combining backup, copy, HSM and device driver routines into a single package. So DFSMS is actually multiple products; it is a suite of data and storage management offerings:-

- DFSMSdftp : Data Facility Product - provides the logical and physical input and output for z/OS storage, it keeps track of all data and programs managed within z/OS, and it provides data access both for native z/OS applications and other platforms.
- DFSMSdss : this is a priced optional feature. It is a DASD data and space management tool for moving and copying data.
- DFSMSHsm : Hierarchical Storage Manager - a priced optional feature for managing low-activity and inactive data. It provides backup, recovery, migration, and space management functions.
- DFSMSRmm : Removable Media Manager - a priced optional feature for managing removable media resources (e.g. IBM's Virtual Tape Server).
- DFSMSStvs : Transactional VSAM Services – is another priced optional feature that enables batch jobs and CICS online transactions to update shared VSAM data sets concurrently.

But we are not going to delve into all of these various components. What we are most interested in is how DB2 can benefit from DFSMS. Using DFSMS, a DB2 DBA can simplify the interaction of DB2 database creation and storage specification. It can deliver:-

- Simplified data allocation
- Improved allocation control
- Improved performance management
- Automated disk space management
- Improved data availability management
- Simplified data movement

DFSMS has the necessary flexibility to support everything DB2 DBAs may want to accomplish in terms of data set placement and design. In this day and age there is no reason to not take advantage of DFSMS for DB2 data sets. However, to achieve a successful implementation, an agreement between the storage administrator and the DB2 administrator is required so that they can together establish an environment that satisfies both their objectives.

SMS Data Classes, Storage Classes, Management Classes, and Storage Groups can be used along with SMS ACS (Automatic Class Selection) routines to set up and automate DB2 data set allocation and placement. SMS Storage Groups contains volumes that satisfy the service requirements of the data sets allocated to them. They can handle more than one type of data. Separate Storage Groups should be defined for production table spaces, active logs, other production data, and non-production data. ACS routines assign data sets to SMS storage classes. For example, indexes can be assigned to one SMS storage class and table spaces to a different SMS storage class. For DB2 Storage Groups using SMS the volume list is set to '*'.

As of DB2 9, DATACLAS, MGMTCLAS, and STORCLAS can be specified in DB2 Storage Groups, and if so, then the VOLUMES clause can be omitted. If the VOLUMES clause is omitted, the volume selection is controlled by SMS. Keep in mind, though, that when processing the VOLUMES, DATACLAS, MGMTCLAS, or STORCLAS clauses, DB2 does not check the existence of the volumes or classes or determine the types of devices that are identified or if SMS is active. Later, when the storage group allocates data sets, the list of volumes is passed in the specified order to Data Facilities (DFSMSdfp).

Basically, using SMS with DB2 is the sane thing to do these days because the new disk architectures, with concepts like log structured files and with cache in the gigabyte sizes, render conventional database design rules based on data set placement less important. In most cases, placement is not an issue, and when it is, SMS classes and ACS routines can be setup to take care of things.

What About Extents?

Some folks think “With RAID/modern storage devices and new DB2 and z/OS features, extents are no longer anything to worry about.” But this is not exactly true. For one thing, the latest extent management features only work with SMS-managed data sets, so if you are still using user-managed data sets then all of the old rules apply!

For SMS-managed data set you can have up to 123 extents on each of 59 volumes. So as of z/OS 1.7, the limit is 7,257 extents for a data set instead of the 255 we’ve been used to for some time. Again though, to enable this requires DFSMS (modify the DFSMS Data Class to set the Extent Constraint Removal to YES).

Extent consolidation also requires SMS-managed STOGROUPS. If a new extent is adjacent to old, they will be merged together automatically. This can result in some extents being larger than the PRIQTY or SECQTY specification(s). Note that this feature was introduced in z/OS 1.5.

OK, so what if everything is SMS-controlled? Then extents don’t matter, right? Well, not really. Even then it is possible for extents to impact performance. Each extent on a disk file has different control blocks controlling access. This means that elapsed time can increase if there is heavy insert activity. For other types of processing (read and update) the number of extents really does not impact on performance.

At any rate, things are not like the olden days where you had to regularly monitor extents and clean them up all the time by reorganizing your table spaces and index spaces. Oh, we want to clean those extents up periodically, but storage administrators have other methods of reducing extents that perhaps can be quicker and/or easier, for example.

- DFSMSHsm MIGRATE and RECALL functions
- DFSMSdss COPY or DUMP and RESTORE functions
- DEFRAg with the CONSOLIDATE keyword
- Other products: e.g. Real Time Defrag

As of V8, DB2 can allocate sliding scale secondary extents. This is enabled by setting MGEXTSZ DSNZ-PARM to YES. Note that the default is NO for V8, but changed to YES automatically when you upgrade to DB2 9. With sliding scale extents the extent sizes allocated gradually increase. DB2 uses a sliding scale for secondary extent allocations of table spaces and indexes when:

- You do not specify a value for the SECQTY option of a CREATE TABLESPACE or CREATE INDEX statement
- You specify a value of -1 for the SECQTY option of an ALTER TABLESPACE or ALTER INDEX statement.

Otherwise, DB2 uses the SECQTY value for secondary extent allocations, if one is explicitly specified (and the SECQTY value is larger than the value that is derived from the sliding scale algorithm). If the table space or index space has a SECQTY greater than 0, the primary space allocation of each subsequent data set is the larger of the SECQTY setting and the value that is derived from a sliding scale algorithm.

Without going into all of the gory details, sliding scale extent allocation can help to reduce the number of extents for your Db2 objects as they grow in size over time. And it can help when you do not have a firm understanding of how your data will grow over time.

Another Thing to Think About

Every now and then I hear from a DB2 DBA complaining about index growth. They'll say something like this: "My index keeps growing and growing and taking additional extents, even after deleting data from the base table. What is going on?"

Well, deleted index keys are not physically deleted, but marked as pseudo-deleted. This can cause an index to grow even as the table data remains at relatively the same level. And it can result in a high CPU cost for index scans because DB2 scans all of the entries, even the pseudo-deleted ones. For this reason, you should keep an eye on tables with a lot of delete activity and periodically reorganize their indexes. You can track pseudo-deleted index keys in SYSINDEXPART if using RUNSTATS or SYSINDEXSPACESTATS if using real time statistics. Consider reorganizing these indexes when pct of pseudo-deleted entries is:

- Greater than 10% for non Data Sharing
- Greater than 5% for Data Sharing

Best Practices

In general, you should adopt best practices for managing your DB2-related storage. Keep up-to-date on DB2/storage functionality and adopt new practices in the latest releases of DB2.

It is a good idea to perform regular and proactive monitoring. Examples of things you should be tracking include:

1. Space display and monitoring for your entire DB2 system.
2. Space display and monitoring of individual DB2 databases.
3. Space display and monitoring of all of your table spaces and indexes.
4. Display and monitoring of the Storage Groups and the associated volumes of a DB2 system. (Data, Workfile, Image Copies, Logs, Archives, Sort/Work etc.)
5. The ability to display all VSAM data sets for all table spaces and indexes (Used, Allocated, Primary and Secondary Quantity, Volumes) and monitoring of the extents for each.
6. Display of the Page Sets of table spaces and indexes that reach their maximum size and maximum number of data sets.
7. Tracking of image copy backup data sets, including the intelligent HSM migration of same. You should be able to reduce backups by track which are not used for local recovery (to CURRENT), as well as data sets older than the last full image copy (including dual and remote backups).
8. Managing the deletion of Image copy backup datasets that are no longer needed because of DROP, DROP/CREATE or MODIFY TABLESPACE ('orphaned', not listed in SYSIBM.SYSCOPY).

If possible, build alerts to inform you of problems, shortages, and potential errors. When possible, automate remediation tactics so that the alert tells you what happened, as well as what was done to correct the issue. Tools may be able to assist in automating reaction to shortages, potential errors, superfluous data sets, etc.

Summary

Finally, there is a people issue that needs to be addressed. The DBA and the storage administrator will need to cooperate in order to facilitate proper database storage. Remember, other types of data that are not stored in the DBMS will need to be saved on disk, too. Databases use storage differently than non-database data. Indexing, partitioning, clustering, and separation of data will cause the database to require more storage, across more drives, and using cache differently than most storage administrators may anticipate.

DBAs need to work storage administrators to make sure that each understands the other domain. And to make sure that s/he has the storage information needed to properly administer DB2. The better the DBA communicates, and the better the relationship is between these two IT professionals, the better your database applications will perform. And that is what it is all really about, right?

Author Bio

Craig S. Mullins is a data management strategist, researcher, and consultant. He is president and principal consultant of Mullins Consulting, Inc., a principal with SoftwareOnZ (a mainframe software distributor), and the publisher/editor for TheDatabaseSite.com. Craig has nearly three decades of experience in all facets of database systems development and has worked with mainframe DB2 since V1.

You may know Craig from his popular books: "DB2 Developer's Guide" (with over 1500 pages of in-depth technical information on DB2 for z/OS) and "Database Administration: The Complete Guide to Practices and Procedures" (the industry's only comprehensive guide to heterogeneous database administration). More information available at www.craigsmullins.com.

Don't Flip Out! How to Stop Your Query Access Plans from Flopping

John Hornibrook

Editors : Guy Lohman, Terry Purcell

Has this ever happened to you? You have an important report query that has been running on an IBM® DB2® for Linux®, UNIX®, and Windows® data server consistently in 25 seconds for the past year. Catalog statistics are regularly collected for tables referenced in the query using the RUNSTATS utility. Suddenly, one day this report query takes 250 seconds to execute after statistics have been collected. Or another query in the same reporting application consistently executes in 15 seconds. A small change is made to one of the search conditions in this query that doesn't significantly increase the size of the result set, and execution now takes 500 seconds! After following standard performance problem determination procedures, you discover that the access plan has changed in both scenarios. Now what do you do?

This edition of the Intelligent Optimizer will discuss some of the reasons why access plans occasionally change for the worse and what you can do to prevent access plan flip-flop from happening.

A Reactive Approach

Once you've determined that a query's performance degradation is due to an access plan change, you can compare the differences to try to understand why the DB2 data server optimizer chose a new access plan. The explain facility is the recommended tool for capturing access plans. The access plan's details are written to a set of explain tables. You can use the db2exfmt tool to format the contents of the explain tables to produce a text file, or you can use the Data Studio Visual Explain tool to provide a graphical view of the access plan. You can read more about the explain facility here [<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.admin.perf.doc/doc/c0005134.html>]

The explain facility provides detailed information about the access plan, including each access plan operator, their sequence of execution, the predicates they apply, the number of rows they process, and their estimated cost. The number of rows processed by an access plan operator directly affects its cost. Sometimes the optimizer mis-estimates the number of rows – or cardinality – because of missing or outdated catalog statistics. The inaccurate cardinality estimates may result in inappropriate access plan changes, so verifying them is often a good place for you to start. The cardinality estimates are based on the number of rows in each referenced database object (table, nickname, table function or materialized query table) and the predicates applied to those objects. It is possible that the cardinality catalog statistic for a base object is inaccurate or it is possible that the selectivity estimate for one or more of the predicates applied to the base object is inaccurate. You can verify the optimizer's cardinality estimates by running subsets of the original query to return a count of the actual number of rows processed by specific access plan operators. Alternatively, if DB2 9.7 is being used, you can enable the DB2 data server to automatically capture the actual runtime cardinality of each operator and store them in the explain tables, where you can later format the contents using db2exfmt1. You can read more about capturing actual runtime cardinalities here [<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.admin.perf.doc/doc/c0056362.html>]

Once you've identified and corrected any inaccurate cardinality estimates - where that is possible - the access plan may still not have changed, or it may have changed but is still not performing as well as the original access plan. You will then need to perform deeper analysis to understand why the optimizer is still not choosing the better access plan.

As you can appreciate by this point, analyzing poorly performing access plans requires considerable time and skill. While access plan analysis skills are very valuable for you to have, performing such a deep analysis for more than a relatively small number of query access plan changes is probably impractical. Quite often, access plan problems need to be corrected quickly, and there isn't time to perform the steps described above. So a more proactive approach is preferable.

A Proactive Approach

Rather than scrambling to fix query access plan problems in panic situations, there are a number of modifications that you can make to your SQL statements, data server configuration, and catalog statistics, to avoid unexpected access plan changes. These modifications allow the optimizer to more accurately cost access plans so that queries run faster. Moreover, when access plan cost estimates are accurate, access plans are less likely to change unexpectedly for the worse, when some aspect of the environment that is considered by the optimizer's cost model changes. For example, a small increase in the number of rows in a table should not result in a new access plan that runs significantly slower. Access plans usually change incorrectly in this type of situation because of other costing inaccuracies. In other words, the optimizer's view of the world is distorted, so it is making decisions based on bad data. The following sections in this article describe how you can improve the optimizer's "vision" so it has a clear view of the best access path.

Writing SQL Statements

SQL is a powerful language that enables you to specify relational expressions in syntactically different but semantically equivalent ways. However, some semantically equivalent variations are easier to optimize than others. Although the DB2 data server optimizer has a powerful query rewrite capability, it might not always be able to rewrite an SQL statement into the most optimal form. There are also certain SQL constructs that can limit the access plans considered by the query optimizer. For example, complex expressions in predicates can limit the optimizer's ability to compute accurate selectivity estimates and can prevent certain access plan operators from being used. The following are some examples of such predicates:-

1. WHERE SALES.PRICE * SALES.DISCOUNT = TRANS.FINAL_PRICE
2. WHERE UPPER(CUST.LASTNAME) = TRANS.NAME
3. WHERE INTEGER(TRANS_DATE)/100 = 200802

All of these predicates may prevent the optimizer from computing an accurate selectivity estimate. The first two predicates will limit the join method to nested-loop join because hash join and merge join cannot be used if the join predicate contains an expression. The third predicate cannot be applied by an index on the TRANS_DATE column using start/stop keys.

One approach you can use to avoid expressions in predicates is to materialize the expression in the table as a generated column. An index can then be created on the column. For example:

```
CREATE TABLE CUSTOMER
(LASTNAME   VARCHAR(100),
 U_LASTNAME VARCHAR(100)
    GENERATED ALWAYS AS (UCASE(LASTNAME))
);

CREATE INDEX CUST_U_LASTNAME ON CUSTOMER(U_LASTNAME)
```

Another approach is to use the inverse of the expression, to avoid referencing a column in the expression. With this approach, the column is referenced alone on one side of the predicate. For example, you can write the third predicate above as:

```
WHERE TRANS_DATE BETWEEN 20080201 AND 20080229
```

This allows the BETWEEN predicate to be applied as a start and stop key on an index on TRANS_DATE.

This section [<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.admin.perf.doc/doc/c0054707.html>] of the DB2 9.7 Information Center contains more information about best practices for writing good SQL statements.

Another potential problem for query optimization is the use of input variables, such as host variables and parameter markers, in predicates, rather than literals. For example:

```
WHERE TRANS_DATE BETWEEN ? AND ?
```

When literals are specified in the predicate, the optimizer can compare them to frequent value or histogram statistics to compute a more accurate selectivity estimate. Using these distribution statistics will result in a more accurate selectivity estimate when the data is skewed or when range predicates are used.

Input variables are essential for good statement preparation time in an online transaction processing (OLTP) environment, where statements tend to be simpler and query access plan selection is more straightforward. Multiple executions of the same query with different input variable values can reuse the compiled access section in the dynamic statement cache, avoiding expensive SQL statement compilations whenever the input values change.

However, for complex queries, input variables can cause problems because plan selection is more complex and the optimizer needs more information to make the best decisions. Moreover, for complex queries, statement compilation time is usually a small component of total execution time, and complex queries do not tend to be repeated with exactly the same predicates, so they do not benefit from the dynamic statement cache.

If input variables need to be used in a complex query workload, consider using the REOPT bind option. The REOPT bind option defers statement compilation from PREPARE to OPEN or EXECUTE time, when the input variable values are known. The values are passed to the optimizer so it can use them to compute a more accurate selectivity estimate. The REOPT bind option can also be used for predicates that reference special registers, such as:

```
WHERE TRANS_DATE < CURRENT DATE
```

This Information Center page [<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.admin.perf.doc/doc/c0055082.html>] contains more details on how to use the REOPT bind option, for different types of applications.

Data Server Configuration

There are a number of configuration options that affect the DB2 data server's performance, such as database and database manager configuration parameters, registry variables, bind options, special registers, and database object attributes. Some of these configuration options have a direct effect on the DB2 data server optimizer because they provide information about the system's capabilities that is considered

when estimating the cost of access plans. Therefore, access plans will tend to be more stable if configuration options important to the optimizer's cost model are set accurately.

Configuration Parameters

The following two database manager configuration parameters are used by the optimizer's cost model. While both are set automatically by the DB2 data server, you should review them to ensure that they are set to appropriate values for your environment.

Table 1.

CPUSPEED	The average time to execute a machine instruction in milliseconds. This value is automatically computed by the DB2 data server. It can also be set manually.
COMM_BANDWIDTH	The communication bandwidth between database partition servers in megabytes per second. This value is automatically computed by the DB2 data server based on the speed of the underlying communications adapter. This value is used by the query optimizer to estimate the cost of transferring data between the database partition servers of a partitioned database system.

The optimizer's cost model also considers the size of the buffer pools and the amount of available sortheap memory. The cost model also considers the amount of memory available for acquiring locks, so that it can choose the best lock granularity appropriate for the specified isolation level. The size of these memory resources should be determined based on available system memory and the type of workload running on your data server. You can specify the size of these memory resources yourself, or you can let the DB2 data server do it for you, using the self-tuning memory feature. Whatever approach you choose, you should be aware that the size of these memory resources is considered by the optimizer too.

The AVG_APPLS database configuration parameter is used by the query optimizer to help estimate how much buffer pool will be available at runtime for the chosen access plan. This parameter is only used by the optimizer's cost model and doesn't directly affect the allocation of system resources at query execution time. When running the DB2 data server in a multi-user environment, particularly with complex queries and a large buffer pool, this parameter informs the query optimizer that multiple query users will be using your system, so that the optimizer should be more conservative in its assumptions of buffer pool availability. However, setting this parameter too high may result in access plans that favor using more CPU or sortheap, due to incorrectly estimating that extra I/O will occur because of insufficient buffer pool memory. It is recommended that this parameter not be set higher than 3 or 5, if it is even necessary to set it at all.

Table Space Attributes

The optimizer also considers storage subsystem characteristics in the cost model. The storage subsystem characteristics are represented as table space attributes. You can specify table space attributes using the CREATE or ALTER TABLESPACE statements. The optimizer considers the following table space attributes.

Table 2.

PREFETCHSIZE	The number of pages to be read when prefetching is performed.
EXTENTSIZE	The size of each extent in pages. This many pages are written to one container in the table space before switching to the next container. This parameter can only be specified when a table space is created.

OVERHEAD	Provides an estimate of the time in milliseconds that is required by a table space container before any data is read into memory. This overhead activity includes the container's I/O controller overhead as well as the disk latency time, which includes the disk seek time.
TRANSFERRATE	Provides an estimate of the time in milliseconds that is required to read one page of data into memory

The DB2 data server doesn't automatically determine the table space overhead and transfer rate. If you don't specify values when the table space is created, default values are used by the optimizer. However they may not be representative of your storage subsystem. I/O can be the most significant cost of an access plan, so it is important for you to set these parameters accurately. This Information Center page [<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.admin.perf.doc/doc/c0005051.html>] contains information on how to set these parameters accurately.

Optimization Class

The DB2 data server allows you to control the level of optimization performed by the query optimizer. Different levels of optimization are grouped into classes. You can specify the query optimization class using the `CURRENT QUERY OPTIMIZATION` special register or the `QUERYOPT` bind option. Setting the optimization class can provide some of the advantages of explicitly specifying optimization techniques, particularly for the following reasons:

- To manage very small databases or very simple dynamic queries
- To accommodate memory limitations at compile time on your database server
- To reduce the query compilation time, such as `PREPARE`.

Most statements can be adequately optimized with a reasonable amount of resources by using optimization class 5, which is the default query optimization class. At a given optimization class, the query compilation time and resource consumption is primarily influenced by the complexity of the query, particularly the number of joins and subqueries. However, compilation time and resource usage are also affected by the amount of optimization performed.

Query optimization classes 1, 2, 3, 5, and 7 are all suitable for general-purpose use. Consider class 0 only if you require further reductions in query compilation time and you know that the SQL statements are extremely simple.

Be aware that if you choose an optimization class that is too low for your workload, the estimated cost of the access plans may be inaccurate because certain types of statistics aren't used. In particular, frequent value and histogram statistics aren't used at optimization class 0 and 1 and histogram statistics aren't used at optimization class 3. This DB2 Information Center page [<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.admin.perf.doc/doc/c0005277.html>] provides more details on the different optimization classes and how to choose the best optimization class for your workload.

Catalog Statistics

The single most important action that you can take to stabilize your access plans is to ensure that accurate catalog statistics are available. You can collect statistics by issuing the `RUNSTATS` command or you can configure the DB2 data server to automatically collect statistics for you. You can also request for statistics to be collected when using the `LOAD`, `CREATE INDEX`, or `REDISTRIBUTE` utilities.

If collecting statistics manually using the RUNSTATS command, you should use the following options at a minimum:

```
RUNSTATS ON TABLE DB2USER.DAILY_SALES WITH DISTRIBUTION  
AND SAMPLED DETAILED INDEXES ALL;
```

Distribution statistics make the optimizer aware of data skew. Detailed index statistics provide more details about the I/O required to fetch data pages when the table is accessed using a particular index. Collecting detailed index statistics consumes considerable CPU and memory for large tables. The SAMPLED option provides detailed index statistics with nearly the same accuracy but requires a fraction of the CPU and memory. These options are also used by automatic statistics collection when a statistical profile has not been provided for a table.

To further improve query performance and access plan stability, you should consider collecting more advanced statistics such as column group cardinalities, LIKE statistics or creating statistical views.

Column Group Statistics

Column group cardinality statistics represent the number of distinct values in a group of columns. This statistic helps the optimizer recognize when the data in columns is correlated. Without it, the optimizer assumes that the data in columns is independent. For example, consider an automobile table containing a MAKE and MODEL column with 20 and 200 distinct values, respectively. Assuming each value of MAKE and MODEL occurs uniformly, the optimizer will compute the combined selectivity estimate of these predicates:

```
WHERE MAKE='HONDA' AND MODEL='ODYSSEY'
```

as $(1/200) * (1/20) = 1/4000$. However, it is unlikely that a particular model of automobile occurs for more than one make, so the number of distinct make and model combinations is likely to be 200 and the combined selectivity of these predicates is $1/200$. You can collect the column group cardinality for these columns using this RUNSTATS command:

```
RUNSTATS ON TABLE DB2USER.AUTOS ON ALL COLUMNS AND COLUMNS((MAKE,MODEL))  
WITH DISTRIBUTION AND SAMPLED DETAILED INDEXES ALL;
```

LIKE Predicate Statistics

If you specify LIKE predicates using the % wildcard character in any position other than at the end of the pattern, you should collect basic information about the sub-element structure of the strings in your data. For example:

```
WHERE KEYWORDS LIKE '%simulation%'
```

Table columns should contain sub-fields or sub-elements separated by blanks. For example, a four-row table DOCUMENTS contains a KEYWORDS column with lists of relevant keywords for text retrieval purposes. The values in KEYWORDS are:

```
'database simulation analytical business intelligence'  
'simulation model fruit fly reproduction temperature'  
'forestry spruce soil erosion rainfall'  
'forest temperature soil precipitation fire'
```

In this example, each column value consists of 5 sub-elements, each of which is a word (the keyword), separated from the others by one blank.

When the % wildcard character appears in some position other than at the end of the pattern, the optimizer assumes that the column being matched contains a series of elements concatenated together, and it estimates the length of each element based on the length of the string, excluding leading and trailing % characters. If you collect sub-element statistics, the optimizer will have information about the length of each sub-element and the delimiter. It can use this additional information to more accurately estimate how many rows will match the predicate.

To collect sub-element statistics, execute RUNSTATS with the LIKE STATISTICS clause.

Statistical Views

The statistics described above are all important for computing an accurate cardinality estimate; however there are some situations where more sophisticated statistics are required. In particular, more sophisticated statistics are required to represent more complex relationships, such as predicates involving expressions:

```
WHERE PRICE > MSRP + MARKUP
```

relationships spanning multiple tables:

```
WHERE PRODUCT.NAME='BBQ' AND PRODUCT.PRODKEY = DAILY_SALES.PRODKEY
```

or anything other than predicates involving independent attributes and simple comparison operations. Statistical views are able to represent these types of complex relationships because statistics are collected on the result set returned by the view, rather than the base tables referenced by the view.

When a query is compiled, the optimizer automatically tries to match the query to the available statistical views. When the optimizer computes cardinality estimates for intermediate result sets, it uses the statistics from matching views to compute a better estimate.

Queries do not need to reference the statistical view directly in order for the optimizer to use it. The optimizer uses the same matching mechanism that is used for materialized query tables (MQTs) to match queries automatically to statistical views. In this respect, statistical views are very similar to MQTs, except they are not stored permanently, so they do not consume disk space, and do not have to be maintained.

However, you do have to decide what statistical views to create, and create them. You create a statistical view by first creating a view and then enabling it for optimization using the ALTER VIEW statement. You then issue RUNSTATS on the statistical view. For example, in order to create a statistical view to represent the join between the time dimension table and the fact table in a star schema, you would do the following:

```
CREATE VIEW SV_TIME_FACT AS  
(SELECT T.* FROM TIME T, DAILY_SALES S  
  WHERE T.TIMEKEY = S.TIMEKEY
```

```
) ;  
  
ALTER VIEW SV_TIME_FACT ENABLE QUERY OPTIMIZATION;  
  
RUNSTATS ON TABLE DB2DBA.SV_TIME_FACT WITH DISTRIBUTION;
```

This statistical view can be used to improve the cardinality estimate - and consequently the access plan and query performance - for queries such as:

```
SELECT SUM(S.PRICE)  
FROM DAILY_SALES S, TIME T, PRODUCT P  
WHERE  
T.TIMEKEY = S.TIMEKEY AND T.YEAR_MON = 200712 AND  
P.PRODKEY = S.PRODKEY AND P.PROD_DESC = 'Power drill';
```

Without the statistical view, the optimizer assumes that all fact table TIMEKEY values corresponding to a particular time dimension YEAR_MON value occur uniformly within the fact table. However, sales might have been particularly strong in December, resulting in many more sales transactions than other months.

There are many situations where statistical views can improve query performance and improve access plan stability. There are some straightforward best practices available to help you determine what statistical views to create. You can find more information about how statistical views work and when to create them in this DB2 Information Center page. [<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.admin.perf.doc/doc/c0021713.html>] .

Handling Exceptions

Unfortunately, occasionally access plans continue to unexpectedly change for the worse, even though you've taken the tuning actions described above. Sometimes it is necessary for you to override the optimizer's decisions, in order to quickly correct a poorly performing access plan or to reduce the risk of an existing access plan changing. If you've properly tuned your SQL statements and your DB2 data server, you should only need to override the optimizer in exceptional situations. Nonetheless, it is important for you to be aware of the different ways to override the optimizer, so that you can react quickly in emergency situations, or so that you can prevent emergency situations from occurring at all.

Access Plan Reuse

Access plan reuse is a new feature available in DB2 9.7 that allows you to request that the access plans for static SQL statements be preserved across binds or rebinds.

In some situations, you might choose to rebind your packages in order to update the access plans to reflect changes in catalog statistics or changes in the DB2 data server configuration. However, sometimes existing packages need to be rebound because of functional changes, such as a release upgrade or the installation of a new service level, but you don't want the existing access plans to change. When a new DB2 data server release is installed, existing packages must be rebound in order for access sections to be rebuilt, so that they are compatible with the runtime component of the new data server release. When a new service level is installed, existing packages don't need to be rebound because access sections are compatible across service levels for the same release. However, you may want to rebuild existing access sections, to take advantage of runtime improvements in the new service level.

Sometimes existing packages need to be rebound because of changes in database objects referenced by SQL statements in the packages. Changes to the database objects may have implicitly invalidated existing packages, so they must be rebound before they can be used. You can rebind invalidated packages yourself using the BIND or REBIND commands or you can let the DB2 data server automatically rebind the packages the next time the application attempts to use the package.

Ideally, access plans should change for the better when a package is rebound, but sometimes you cannot take that risk, because you aren't able to react fast enough to unexpected access plan changes.

You can enable access plan reuse through the ALTER PACKAGE statement, or by using the APREUSE option on the BIND, REBIND, or PRECOMPILE command. Packages that are subject to access plan reuse have the value 'Y' in the APREUSE column of the SYSCAT.PACKAGES catalog view.

The ALTER_ROUTINE_PACKAGE procedure is a convenient way for you to enable access plan reuse for compiled SQL objects, such as routines, functions, and triggers.

Access plan reuse is most effective when changes to the schema and compilation environment are kept to a minimum. If significant changes are made, it might not be possible to recreate the previous access plan. Examples of such significant changes include dropping an index that is being used in an access plan, or recompiling an SQL statement at a different optimization level. Significant changes to the query compiler's analysis of the statement can also result in the previous access plan no longer being reusable.

Access plans from packages that were produced by releases prior to DB2 9.7 cannot be reused.

If an access plan cannot be reused, compilation continues, but a warning (SQL20516W) is returned with a reason code that indicates why the attempt to reuse the access plan was not successful. Additional information is sometimes provided in the diagnostic messages that are available through the explain facility.

You can read more about access plan reuse in this DB2 Information Center page [<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.admin.perf.doc/doc/c0055383.html>].

Optimization Profiles

The DB2 data server supports an even more direct way to influence access plans using optimization profiles. Optimization profiles allow you to specify access plan details such as base access and join methods and join order. For example, you can specify that access to a particular table should use a particular index or you can specify that two tables should be joined using the hash join method. Optimization profiles also allow you to control query rewrite optimizations such transforming certain types of subqueries to joins. You can specify the base table access methods, join methods, and join order for the entire access plan, or just a subset of the access plan.

An optimization profile is specified as an XML document that you create and store in the SYSTOOLS.OPT_PROFILE table. You put an optimization profile into effect for your application by specifying the optimization profile name using the CURRENT OPTIMIZATION PROFILE special register for dynamic SQL statements, or by specifying the OPTPROFILE bind option for static SQL statements.

An optimization profile contains optimization guidelines that specify the access plan details. An optimization profile can contain optimization guidelines for one or more SQL statements. The SQL statement text is stored in the optimization profile along with the optimization guidelines. When an optimization profile is in effect for your application, each SQL statement compiled by your application will be matched to the SQL statements specified in the optimization profile. When a matching SQL statement is found in the optimization profile, the SQL compiler will use the optimization guidelines for that SQL statement while optimizing it.

Example 1. An example of an optimization profile

```
<?xml version="1.0" encoding="UTF-8"?>
<OPTPROFILE VERSION="9.7.0">
  <!--
    Global optimization guidelines section. Optional but at most one.
  -->
  <OPTGUIDELINES>
    <RTS TIME="1000"/>
  </OPTGUIDELINES>
  <!--
    Statement profile section. Zero or more.
  -->
  <STMTPROFILE ID="Guidelines for Q1">
    <STMTKEY SCHEMA="DB2USER">
      <![CDATA[SELECT SUM(S.PRICE)
        FROM DAILY_SALES S, TIME T, PRODUCT P
        WHERE
          T.TIMEKEY = S.TIMEKEY AND T.YEAR_MON = 200712 AND
          P.PRODKEY = S.PRODKEY AND P.PROD_DESC = 'Power drill']]>
    </STMTKEY>
    <OPTGUIDELINES>
      <HSJOIN>
        <NLJOIN>
          <IXSCAN TABLE="T" INDEX="TD_IX1"/>
          <IXSCAN TABLE="S" INDEX="DS_TIMEKEY"/>
        </NLJOIN>
        <ACCESS TABLE="P" />
      </HSJOIN>
    </OPTGUIDELINES>
  </STMTPROFILE>
</OPTPROFILE>
```

This optimization profile specifies an optimization guideline for one SQL statement. The optimization guideline is represented by the OPTGUIDELINES element. This particular optimization guideline specifies that the TIME table should be accessed with the TD_IX1 index and then joined using a nested-loop join to the DAILY_SALES table which will be accessed on the inner of the join using index DS_TIMEKEY. The result of this join will be joined to the PRODUCT table using a hash join. The PRODUCT table will be on the inner of the hash join. The access to the PRODUCT table is specified using the ACCESS element, which means that the optimizer is free to choose the access method.

Tables are referenced in the optimization guideline by using the TABLE element to specify the correlation name used to reference the tables in the original SQL statement. For example, the DAILY_SALES table is given the correlation name S in the original SQL statement, so you must reference it as S when using the TABLE element.

You may also specify global optimization guidelines in the optimization profile. Global optimization guidelines specify some aspect of the SQL statement compilation environment, such as the optimization class, degree of parallelism, or eligible MQTs. Global optimization guidelines are applied to all SQL statements that are compiled while the containing optimization profile is in effect. The global optimization guideline in this example specifies a time limit of 1000 ms for collecting statistics at SQL statement compilation time, by the real-time statistics collection feature.

Optimization profiles are a powerful tool for controlling access plans; however, they should be used with caution. Optimization profiles prevent access plans from adjusting to changes in your data and your environment. While this does accomplish access plan stabilization, it may be a bad approach when used for extended periods of time, because the performance improvements resulting from better access plans will never be realized. Optimization profiles are best used for exceptional situations when the tuning actions described above are unsuccessful in improving or stabilizing access plans.

This DB2 Information Center page [<http://publib.boulder.ibm.com/infocenter/db2luw/v9r7/topic/com.ibm.db2.luw.admin.perf.doc/doc/c0024522.html>] provides more details on how to use optimization profiles. .

Conclusion

The DB2 data server query optimizer is more likely to choose optimal access plans consistently if you provide it accurate information about your data and system configuration. Access plans will also tend to be more stable because they will be less sensitive to small changes in your data and system configuration. How you write your SQL statements also affects the optimizer's ability to choose optimal access plans. Following the tuning actions described in this article should help you achieve consistently good query performance. However, for the exceptional situations in which the described tuning actions don't have the desired effect, you can override the optimizer using access plan reuse or optimization profiles. The overall combination of system tuning actions and selective optimizer overrides should help you achieve your DB2 data server query performance objectives.

IBM and DB2 are registered trademarks of International Business Machines Corporation in the United States, other countries, or both.

Linux is a registered trademark of Linus Torvalds in the United States, other countries, or both.

UNIX is a registered trademark of The Open Group in the United States and other countries.

Windows is a trademark of Microsoft Corporation in the United States, other countries, or both.

Other company, product, or service names may be trademarks or service marks of others.